

# AI535: Deep Learning

## Assignment 2: CIFAR Image Classification using Fully-Connected Network

Name: JuHyun Kim

### 1) Implementation of Layers

Linear Layer: Stores weights, biases, and their velocities for momentum updates. The forward pass computes  $XW + b$ , and the backward pass computes gradients for both inputs and parameters.

```
class LinearLayer:
    # TODO: Initialize our layer with (input_dim, output_dim) weight matrix and a (1,output_dim) bias vector
    def __init__(self, input_dim, output_dim):
        # He initialization for weights
        self.weights = np.random.randn(input_dim, output_dim) * np.sqrt(2. / input_dim)
        self.bias = np.zeros((1, output_dim))
        self.velocity_weights = np.zeros_like(self.weights)
        self.velocity_bias = np.zeros_like(self.bias)
    # TODO: During the forward pass, we simply compute  $XW+b$ 
    def forward(self, input):
        self.input = input
        output = np.dot(input, self.weights) + self.bias
        return output
    def backward(self, grad):
        self.grad_weights = np.dot(self.input.T, grad)
        self.grad_bias = np.sum(grad, axis=0, keepdims=True)
        grad_input = np.dot(grad, self.weights.T)
        return grad_input
```

ReLU Layer:

On this function, stores the input to determine where to apply gradients during the backward pass, allowing gradients to flow only through the activations that were positive in the forward pass.

```
class ReLU:
    def forward(self, input):
        self.input = input
        return np.maximum(0, input)
    def backward(self, grad):
        return grad * (self.input > 0).astype(grad.dtype)
    # No parameters so nothing to do during a gradient descent step
    def step(self, step_size, momentum=0, weight_decay=0):
        return
```

Sigmoid Cross Entropy Layer: Combines the sigmoid activation with cross-entropy loss for binary classification. It calculates the loss during the forward pass and the gradient with respect to the input during the backward pass.

```
class SigmoidCrossEntropy:
    def forward(self, logits, labels):
        self.sigmoid_output = 1 / (1 + np.exp(-logits))
        loss = -np.mean(labels * np.log(self.sigmoid_output + 1e-7) + (1 - labels) * np.log(1 - self.sigmoid_output + 1e-7))
        return loss
    def backward(self, logits, labels):
        grad = self.sigmoid_output - labels
        return grad
```

## 2) Mini-Batch Processing

**Forward Pass:** The model's forward method is called with the current mini-batch of inputs ( $X_{\text{batch}}$ ), producing logits (the raw output scores from the model, before applying the sigmoid function).

**Compute Loss:** The SigmoidCrossEntropy loss between the logits and the actual labels ( $Y_{\text{batch}}$ ) is computed. This loss quantifies how well the model's predictions match the true labels.

**Backward Pass:** The gradient of the loss with respect to the logits is computed using the backward method of the loss function. This gradient is then propagated backward through the network (via the `net.backward` method) to compute gradients for all model parameters.

**Optimizer Step:** The model parameters are updated in the direction that minimizes the loss, using stochastic gradient descent (SGD) with momentum and weight decay. This step adjusts the model parameters to reduce the loss on the current mini-batch.

**Accuracy Calculation:** The model's predictions are compared to the true labels to calculate the accuracy for the current mini-batch. The predictions are obtained by thresholding the logits at 0.5, and accuracy is computed as the mean of correct predictions.

**Book-keeping:** The loss and accuracy for each mini-batch are accumulated to compute the average loss and accuracy for the entire epoch.

```
# Q2 TODO: For each epoch below max epochs
for epoch in range(max_epochs):
    epoch_avg_loss = 0.0
    epoch_total_acc = 0.0
    # Shuffle the training data
    permutation = np.random.permutation(num_examples)
    X_train_shuffled = X_train[permutation]
    Y_train_shuffled = Y_train[permutation]
    for i in range(0, num_examples, batch_size):
        # Get mini-batch
        X_batch = X_train_shuffled[i:i + batch_size]
        Y_batch = Y_train_shuffled[i:i + batch_size]
        # Forward pass
        logits = net.forward(X_batch)
        # Compute loss
        loss_func = SigmoidCrossEntropy()
        loss = loss_func.forward(logits, Y_batch)
        epoch_avg_loss += loss
        # Backward pass
        grad_loss = loss_func.backward(logits, Y_batch)
        net.backward(grad_loss)
        # Take optimizer step with momentum
        net.step(step_size, momentum, weight_decay)
        # Book-keeping for accuracy
        predictions = (logits >= 0.5).astype(int)
        accuracy = np.mean(predictions == Y_batch)
        epoch_total_acc += accuracy
    # Calculate average loss and accuracy for the epoch
    epoch_avg_loss /= (num_examples // batch_size)
    epoch_avg_acc = epoch_total_acc / (num_examples // batch_size)
    # Evaluate performance on test set after each epoch
    val_loss, val_acc = evaluate(net, X_test, Y_test, batch_size)
    val_accs.append(val_acc)
    # Add loss and accuracy to lists
    losses.append(epoch_avg_loss)
    accs.append(epoch_avg_acc)
    val_losses.append(val_loss)
```

## 3) Accuracy

Tuned Batch size = 64, learning rate = 0.0001, and hidden layer = 100.

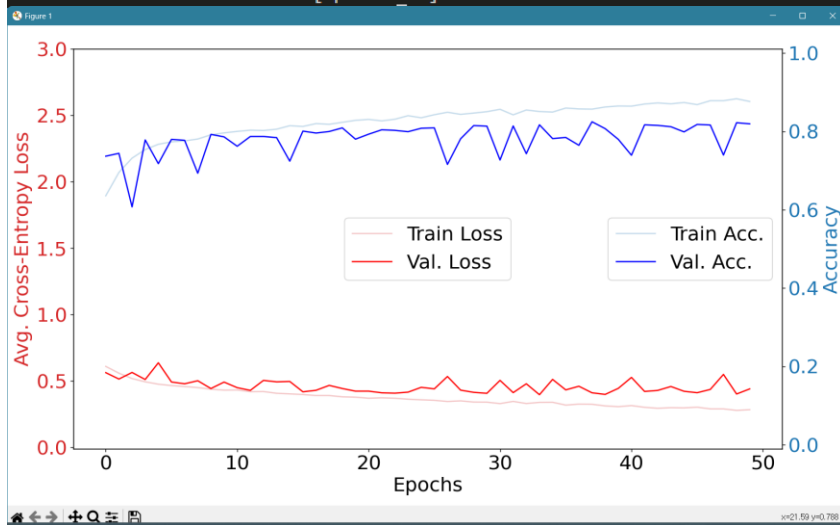
I got 0.248 of Loss, 87.64% of train accuracy and 81.93% of value accuracy at final epoch.

```
def main():
    # Set optimization parameters
    batch_size = 64          #64,128,256
    max_epochs = 50

    step_size = 0.0001      #learning rate 0.01, 0.001, 0.0001
    number_of_layers = 2
    width_of_layers = 100   #hidden layter 50,100,200
    weight_decay = 0.0001
    momentum = 0.8

    # Load data
    data = pickle.load(open('cifar_2class_py3.p', 'rb'))
    X_train = data['train_data'] / 255.0
    Y_train = data['train_labels']
    X_test = data['test_data'] / 255.0
    Y_test = data['test_labels']
```

```
[Epoch 29] Loss: 0.3422 Train Acc: 0.8469 Valid Acc: 0.8149
[Epoch 30] Loss: 0.3415 Train Acc: 0.8504 Valid Acc: 0.8135
[Epoch 31] Loss: 0.3306 Train Acc: 0.8565 Valid Acc: 0.7271
[Epoch 32] Loss: 0.3472 Train Acc: 0.8423 Valid Acc: 0.8140
[Epoch 33] Loss: 0.3311 Train Acc: 0.8547 Valid Acc: 0.7432
[Epoch 34] Loss: 0.3400 Train Acc: 0.8508 Valid Acc: 0.8169
[Epoch 35] Loss: 0.3406 Train Acc: 0.8494 Valid Acc: 0.7817
[Epoch 36] Loss: 0.3194 Train Acc: 0.8597 Valid Acc: 0.7847
[Epoch 37] Loss: 0.3267 Train Acc: 0.8575 Valid Acc: 0.7646
[Epoch 38] Loss: 0.3255 Train Acc: 0.8569 Valid Acc: 0.8247
[Epoch 39] Loss: 0.3134 Train Acc: 0.8624 Valid Acc: 0.8071
[Epoch 40] Loss: 0.3075 Train Acc: 0.8649 Valid Acc: 0.7798
[Epoch 41] Loss: 0.3158 Train Acc: 0.8645 Valid Acc: 0.7393
[Epoch 42] Loss: 0.3027 Train Acc: 0.8700 Valid Acc: 0.8169
[Epoch 43] Loss: 0.2956 Train Acc: 0.8727 Valid Acc: 0.8154
[Epoch 44] Loss: 0.2999 Train Acc: 0.8705 Valid Acc: 0.8120
[Epoch 45] Loss: 0.2985 Train Acc: 0.8738 Valid Acc: 0.7988
[Epoch 46] Loss: 0.3036 Train Acc: 0.8684 Valid Acc: 0.8179
[Epoch 47] Loss: 0.2908 Train Acc: 0.8784 Valid Acc: 0.8164
[Epoch 48] Loss: 0.2909 Train Acc: 0.8784 Valid Acc: 0.7397
[Epoch 49] Loss: 0.2796 Train Acc: 0.8834 Valid Acc: 0.8223
[Epoch 50] Loss: 0.2847 Train Acc: 0.8764 Valid Acc: 0.8193
2024-02-22 22:24:53 INFO [Epoch 9984] Loss: 0.2847 Train Acc: 0.8764% Val Acc: 81.93%
```



#### 4) Training Monitoring – evaluate function

```
#####
# Q4 Implement Evaluation for Monitoring Training
#####

# TODO: Given a model, X/Y dataset, and batch size, return the average cross-entropy loss and accuracy over the set
def evaluate(model, X_val, Y_val, batch_size):
    losses = []
    accuracies = []

    num_examples = X_val.shape[0]

    for i in range(0, num_examples, batch_size):
        X_batch = X_val[i:i+batch_size]
        Y_batch = Y_val[i:i+batch_size]

        logits = model.forward(X_batch)

        loss_func = SigmoidCrossEntropy()
        loss = loss_func.forward(logits, Y_batch)
        losses.append(loss)

        predictions = (logits >= 0.5).astype(int)
        accuracy = np.mean(predictions == Y_batch)
        accuracies.append(accuracy)

    avg_loss = np.mean(losses)
    avg_accuracy = np.mean(accuracies)


    return avg_loss, avg_accuracy
```

For each mini-batch, the `X_batch` and `Y_batch` are the subset of the input data and the corresponding labels are selected to form the mini-batch. And do loss Calculation which the `SigmoidCrossEntropy` object is instantiated, and its forward method is called with logits and `Y_batch` to compute the cross-entropy loss for the mini-batch. This loss is then added to the losses list. The accuracy for the mini-batch is computed as the mean of correct predictions (`predictions == Y_batch`) and is added to the accuracies list. After iterating through all mini-batches, the function calculates the average loss and accuracy across the entire dataset by computing the mean of the losses and accuracies lists, respectively. Finally the function returns `avg_loss` and `avg_accuracy`.

## 5) Tuning Parameters

### i) different number of batch size

Batch size = 256. Increased the batch size, but there are more loss and the accuracy values got lower.



### ii) different learning rate

Learning rate = 0.01. Since the learning rate tuned higher, it can cause the model to overshoot the minimum of the loss function, leading to divergent behavior and instability in training.

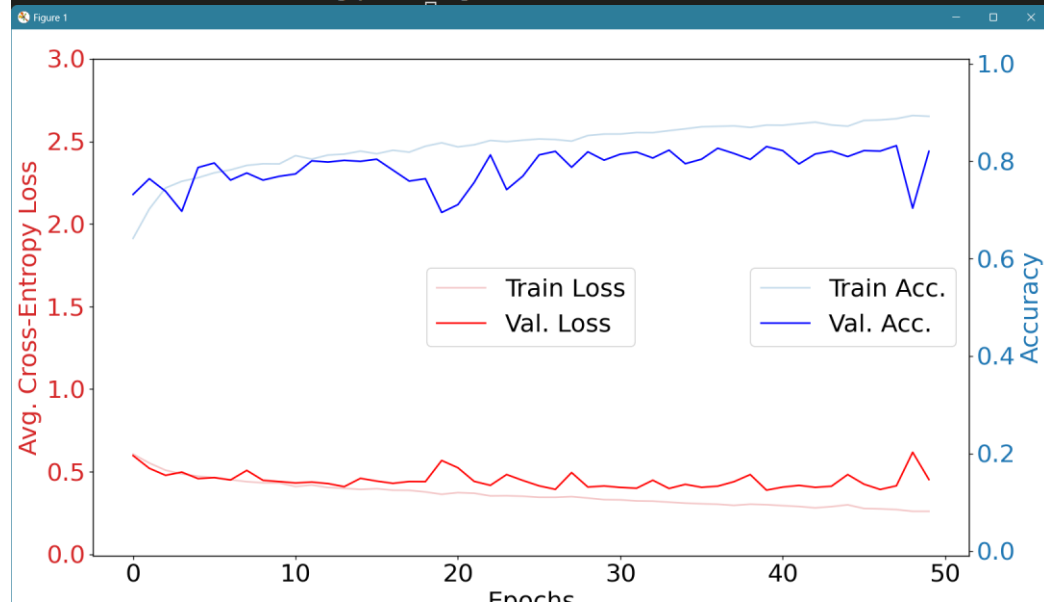
```
[Epoch 39] Loss: 0.7139 Train Acc: 0.5173 Valid Acc: 0.5002
[Epoch 40] Loss: 0.7134 Train Acc: 0.5158 Valid Acc: 0.5002
[Epoch 41] Loss: 0.7197 Train Acc: 0.5159 Valid Acc: 0.5002
[Epoch 42] Loss: 0.7120 Train Acc: 0.5158 Valid Acc: 0.5002
[Epoch 43] Loss: 0.7155 Train Acc: 0.5128 Valid Acc: 0.5002
[Epoch 44] Loss: 0.7143 Train Acc: 0.5128 Valid Acc: 0.5002
[Epoch 45] Loss: 0.7136 Train Acc: 0.5068 Valid Acc: 0.5002
[Epoch 46] Loss: 0.7146 Train Acc: 0.5128 Valid Acc: 0.5002
[Epoch 47] Loss: 0.7136 Train Acc: 0.5053 Valid Acc: 0.5002
[Epoch 48] Loss: 0.7155 Train Acc: 0.5143 Valid Acc: 0.5002
[Epoch 49] Loss: 0.7149 Train Acc: 0.5143 Valid Acc: 0.5002
[Epoch 50] Loss: 0.7164 Train Acc: 0.5113 Valid Acc: 0.5002
2024-02-22 22:34:08 INFO [Epoch 9984] Loss: 0.7164 Train Acc: 0.5113% Val Acc: 50.02%
```



### iii) different number of hidden units

Hidden layers = 200. Adding 100 more hidden layers, therefore the output of training value got higher than when the number of hidden layers is 100. However, there is also has overfitting risk when we increase the data size. But still, this tuned model was the highest achievement I got.

[Epoch 39]	Loss: 0.3031	Train Acc: 0.8694	Valid Acc: 0.8037
[Epoch 40]	Loss: 0.3000	Train Acc: 0.8742	Valid Acc: 0.8301
[Epoch 41]	Loss: 0.2943	Train Acc: 0.8739	Valid Acc: 0.8218
[Epoch 42]	Loss: 0.2892	Train Acc: 0.8772	Valid Acc: 0.7944
[Epoch 43]	Loss: 0.2808	Train Acc: 0.8802	Valid Acc: 0.8149
[Epoch 44]	Loss: 0.2884	Train Acc: 0.8744	Valid Acc: 0.8208
[Epoch 45]	Loss: 0.2993	Train Acc: 0.8719	Valid Acc: 0.8096
[Epoch 46]	Loss: 0.2772	Train Acc: 0.8837	Valid Acc: 0.8218
[Epoch 47]	Loss: 0.2748	Train Acc: 0.8847	Valid Acc: 0.8208
[Epoch 48]	Loss: 0.2705	Train Acc: 0.8873	Valid Acc: 0.8320
[Epoch 49]	Loss: 0.2600	Train Acc: 0.8938	Valid Acc: 0.7036
[Epoch 50]	Loss: 0.2599	Train Acc: 0.8922	Valid Acc: 0.8203
2024-02-22 22:39:20 INFO [Epoch 9984] Loss: 0.2599 Train Acc: 0.8922% Val Acc: 82.03%			



## 6) Discussion

More things to improve, even if I got enough output, the overall performance could be significantly improved by transitioning to architectures more suited for image data, and by incorporating advanced training, regularization, and evaluation techniques. On the implementation of SGD with momentum, the performance of the network could be further enhanced by adopting more sophisticated optimization algorithms like Adam or RMSprop, which adapt the learning rates for each parameter, offering a more nuanced control over the training process. And the evaluation metrics implemented provide a basic understanding of the model's performance. However, relying solely on accuracy and cross-entropy loss may not always offer a complete picture, especially in cases where the dataset might be imbalanced.